

Distinguish Sense from Nonsense: Out-of-Scope Detection for Virtual Assistants

Cheng Qian[†], Haode Qi[†], Gengyu Wang[†], Ladislav Kunc[†], Saloni Potdar[‡]

[†]IBM Watson, [‡]Apple Inc.

{cheng.qian, haode.qi, gengyu, lada}@ibm.com, s_potdar@apple.com

Abstract

Out of Scope (OOS) detection in Conversational AI solutions enables a chatbot to handle a conversation gracefully when it is unable to make sense of the end-user query. Accurately tagging a query as out-of-domain is particularly hard in scenarios when the chatbot is not equipped to handle a topic which has semantic overlap with an existing topic it is trained on. We propose a simple yet effective OOS detection method that outperforms standard OOS detection methods in a real-world deployment of virtual assistants. We discuss the various design and deployment considerations for a cloud platform solution to train virtual assistants and deploy them at scale. Additionally, we propose a collection of datasets that replicates real-world scenarios and show comprehensive results in various settings using both offline and online evaluation metrics.

1 Introduction

In the context of task-oriented dialog, Out of Scope (OOS) detection is the problem of identifying end-user queries that are beyond the scope of a chatbot. While this problem is generally studied under the umbrella of “out of domain” detection in machine learning, we show that unique challenges arise in real-world applications. We study this problem in the context of our enterprise virtual assistant (VA) platform which is used by 10,000+ customers to design chatbots. In this setting, the natural language understanding models comprising of In Scope (IS) and OOS detection modules, need to determine whether an input query belongs to a set of pre-defined intents or if it is out of scope for the chatbot.

Real-world success of OOS systems often involves measuring how good they are at *containment*, i.e., the user queries are resolved and contained by the chatbot while minimizing human interventions. Since containment rate can be only observed after launching the VA online, offline

metrics such as IS accuracy and OOS accuracy are needed while designing and developing the models.

The average designer using an enterprise VA platform is not a machine learning expert. This leads to a variety of challenges in the provided user data, which constitutes the need for robust algorithms. Firstly, end-users often provide data that is heavily imbalanced or noisy for both IS and OOS detection.

While designing VA for enterprise use-cases, IS and OOS examples often naturally belong to the same domain. Such OOS samples are called In Domain OOS (ID-OOS) as opposed to Out-of-Domain OOS (OOD-OOS) which are relatively easier OOS samples from a different domain (Zhang et al. (2022)). Designers expect the VA to detect these relevant but unsupported topics (ID-OOS) even though it has high semantic overlap with IS examples. Finally, while entities defined by the designer play an important role for a real-world VA, they are often ignored in academic OOS settings. We show that entities must be modeled in conjunction with IS and OOS classification.

In this paper, we discuss the challenges of designing a real-world OOS detection system in depth and common approaches taken to design such a system. We propose a simple but effective algorithmic modification for OOS detection in a real-world deployed system. This system models entities, intent and OOS classification jointly and addresses the challenges around data. We propose a comprehensive benchmark based on public datasets and show that our method outperforms standard approaches while being simple to deploy and maintain.

2 Challenges

2.1 Metrics

Containment and Disambiguation For businesses, the key performance index (KPI) metric is typically different from the common machine learning metrics used to test the performance of

the algorithm. Businesses use containment rate to measure chatbot performance - the portion of conversations not handed off or escalated to a human agent for quality reasons. Among offline evaluation metrics, IS performance provides the best estimate for containment rate. Disambiguation is a mechanism to increase containment by asking end-users clarification questions and providing more than one relevant intent. This has to be counter-balanced with high OOS performance so that we don't provide a set of predictions in the form of IS classes for an OOS query. This is essential to appear "intelligent" and handle conversations gracefully.

In Domain Out of Scope Detection refers to detecting OOS samples with high semantic overlap with IS examples in the same domain (ID-OOS). ID-OOS queries are often harder to detect than the easier Out of Domain OOS (OD-OOS) samples. The algorithm should be able to identify ID-OOS and also generalize well to unseen OOD-OOS.

2.2 Data Considerations

The average designer of an enterprise VA platform doesn't need to have ML background, hence the expectations of labeled data are very different from an ML expert end-user.

Class Imbalance is often extreme in data provided by VA designer, with some intent classes having more number of examples than others.

Data-scarce scenarios Labeling data is often expensive for enterprises who desire good performance with very few labelled examples per class, and often no OOS labeled data.

Noisy data Unlike public datasets, enterprise datasets have semantically similar intents due to overlap in business use cases. Additionally, real-world end-user input queries to VAs usually contain spelling errors, intentionally repeated characters, emojis, and slang. Proper normalization is required to improve robustness of OOS algorithms.

2.3 Computational Efficiency

While developing the OOS detection algorithm for an enterprise VA platform, we need to strike a good trade-off between cost of serving the model and performance of the model. Based on our experience, VA platforms are expected to handle training sets of more than 10k training examples and more than 1000 classes.

Model size & memory: There are over 100,000 customer-specific models deployed in production and each chatbot serves millions of queries per

month. Hence low maintenance, training and inference costs can increase profitability.

Training time: Designers typically make changes in an iterative fashion, designing the VA through trial and error. For an interactive experience in the product, the OOS detection component needs to train in 1 minute (Qi et al. (2020)).

Inference time: During the inference, each query passes through all the natural language understanding (NLU) components - IS classification, OOS detection, entity recognition and spellchecking, and needs to provide the predictions in 10 milliseconds.

2.4 Entities and OOS Detection

Entities are designed to represent nouns from end-user inputs and are crucial for VAs to respond accordingly and haven't been studied extensively in OOS detection.

Terminologies Designers can define entities with special terminologies that are out of vocabulary of any other public or private corpus. This requires OOS detection methods to differentiate such terminologies from gibberish sentences.

Synonyms The OOS detection algorithm is expected to produce similar detection scores across the multitudes of synonyms of the same entity.

Numeric Values System entities like date, number, time etc. are pre-configured in a VA to cover a wide range of concepts. However, there is no one-size-fits-all solutions for system entities. E.g., the system entity "11" can be a part of domain specific terminology "operating system Windows 11". The OOS detection algorithm needs to be aware of these system entity values and decide the relevance of the sentence based on the context.

We introduce several potential solutions for handling entities in OOS detection and analyze their advantages and disadvantages.

Concatenation of all entity synonyms In the context of Binary OOS detection, we add one synthetic IS example in to the training data by concatenating all entity synonyms provided for a chatbot. Context independent features such as uni-grams, bi-grams and mean/max pooling of word-embeddings will help recognizing these entities as IS at the runtime. This simple approach works well empirically but has the disadvantage of ignoring the context and semantic meaning.

Synonyms and Entity proxies in intent templates In enterprise VA, an entity can be defined with multiple synonyms. In our product, we support entity proxies, which is a definition of a certain entity and its associated synonyms that are considered

equal. This greatly simplifies training data definition at the cost of potential instability at runtime: our intent detection and OOS algorithm should return the same confidence and same predicted label if one synonym is replaced by another. For the example in Table 3, if "cell phone" is defined as an entity proxy, VA designer only references the symbol "cell phone" in training examples, and at runtime "i want an iphone 11" gets the exact same prediction as "i want an iphone xr".

3 OOS Detection Algorithms

OOS detection algorithms can be broadly classified into single-stage and multi-stage.

3.1 Single-stage OOS

All the IS classes and optionally the OOS class are used together train a single model to determine if the incoming query belongs to one of the IS intents or is OOS.

Multiclass Classification In this approach, the algorithm treats the OOS examples as an additional class as explored in (Zhan et al. (2021), Choi et al. (2021)), alongside the IS classes to train a multiclass classification model for both IS intent detection and OOS detection.¹ This approach trains a single algorithm for both OOS detection and IS detection tasks. In practice, this approach is susceptible to over-fitting to the provided OOS examples and might not generalize well to unseen OOS queries. Additionally, it can fail in the presence of severe class-imbalance.

In-scope Classification utilizing output distribution This type of methods trains a classifier on IS data which outputs a probability vector with low maximum probability or high entropy for an OOS input, as explored in Lewis and Gale (1994), Hendrycks and Gimpel (2016), Lee et al. (2018a), Yilmaz and Toraman (2022). These methods train a single model for IS detection and applies a threshold on output probability distribution statistics (such as max and entropy) for OOS detection. However, in practice, training data typically has semantically overlapped intents which will mislead the system and increase unnecessary human agent intervention as shown in Table 1.

3.2 Multi-stage OOS

Multi-stage OOS method uses a binary classifier to determine if a query is IS or OOS in the first stage.

¹<https://docs.microsoft.com/en-us/azure/cognitive-services/luis/concepts/intents#none-intent>

In the subsequent stages we determine which of the IS intent is the closest match.

Binary Classification (IS/OOS): A binary classifier is trained using the IS examples and OOS examples as explored in Tax and Duin (1999). The classification result is used to determine if end-user query is OOS or ID. In case there are no OOS training examples, the binary OOS classifier can be replaced with an one-class classifier or other unsupervised methods. Another solution for the lack of OOS training data is synthetic OOS training examples, refer to Section 4 for more discussion.

In-scope Classification plus unsupervised methods on internal (hidden state) representation trains a classifier based on IS training examples, and utilizes internal representation (for example, concatenation of hidden states from several layers of a neural network) of the IS classifier for an unsupervised OOS detection algorithm, like autoencoder with reconstruction loss, distance based approach (Wu et al., 2022), (Shen et al., 2021), and density based approach (Lin and Xu, 2019).

3.3 Our Approach

We show a simple modification to the multi-stage OOS to improve the performance of the system and alleviate problems with the other approaches mentioned previously. Our approach **Binary Classification(In Scope/OOS) discounting on In Scope scores** treats OOS classification as a binary classification problem like the previous formulation. However, the binary classification score of the OOS detection algorithm is used to discount the IS classification score to determine the final IS score (more details to follow). In case no training OOS examples are available our OOS detection algorithm becomes one-class classification. This formulation is related to calibration (Kamath et al., 2020) that trains a new model to reject inputs when the model is over-confident. However, our approach applies the OOS output as discounting factor instead of binary score leading to better performance in the context of enterprise VA as shown in Table 4.

In terms of the OOS classification component, we implement a distance based approach based on sentence embedding of both IS and OOS training examples (if labeled). At training time, we first apply the trick described in Section 2.4 to preprocess entities among other text normalization steps, then query the sentence embeddings from a sentence encoder. For each IS example, we store the linear combination of its sentence embedding and the mean embedding of its corresponding intent

class in an approximate nearest neighbor(ANN) search index. If there are OOS training examples, we store their sentence embeddings in the same ANN search index. At runtime, a query is pre-processed the same way as training examples, the cosine distance to the nearest neighbor will be used as OOS score to discount the output from the IS classifier. If the nearest neighbor is OOS, we add an additional constant to the corresponding nearest distance, to discount the confidence more. The discounting step uses the OOS score cos_dist and IS classifier output confidence **conf**, we apply the formula below to produce the final output confidence vector **final_conf** as follows:

$$\mathbf{final_conf} = (1 - f(\max(cos_dist, 0))) \cdot \mathbf{conf} \quad (1)$$

$$f(x) = \begin{cases} x & x \geq 0.5 \\ sigmoid(a \cdot (x - 0.5)) & \text{otherwise} \end{cases} \quad (2)$$

, where a is a constant that can be tuned for different applications. The motivation for Formula 2 is to reduce the amount of discounting on IS confidence (comparing to a linear discounting function), when OOS classifier predicts a low cosine distance (thus high similarity) for an utterance.

Typically, a fixed threshold T on the final output confidences is used in real world applications to determine whether an input utterance is predicted as IS or OOS: a new input is deemed OOS if its final output confidence is less than T . Theoretically, this threshold is not critical to machine learning metrics, especially threshold independent metrics. Even for threshold dependent metrics, this threshold can always be tuned in accordance with the scale of final output confidence to achieve the same results. However in practice, as a commercial VA platform, a fixed threshold reduces the maintenance cost of a chatbot and only a small fraction of the chatbot designers will try to tune the threshold. In our product, 0.2 threshold is set as the default value.

3.4 Benchmark Dataset

We collect 8 intent classification datasets to comprehensively evaluate the methods mentioned above regarding the challenges, including IS classification, OOS detection, and scalability. The 8 datasets include **ATIS** (Hemphill et al., 1990), **BANKING77** (Casanueva et al., 2020), **CLINC150** (Larson et al., 2019), **StackOverflow**², **SNIPS** (Coucke et al., 2018), **HAR** (Liu et al., 2019), **ROSTD** (Gangal et al., 2020), and **HINT3** (Arora et al., 2020). A

²https://storage.googleapis.com/download.tensorflow.org/data/stack_overflow_16k.tar.gz

Query	Intent
I need assistance with my retirement account	retirement account
I need to talk to a agent about my retirement account	agent

Table 1: The two queries shown are semantically overlapped. For the query "I need to talk to a assistant about my retirement account", the correct intent should be "agent" but one can expect "retirement account" and "agent" having similar probability. For approaches that rely on probability vector to detect OOS input, these examples can mislead them to treat valid IS queries as OOS.

summary of dataset statistics after preprocessing is provided in Table 2.

To evaluate methods' performance on ID-OOS detection, we ensure all datasets contain ID-OOS examples. For datasets that only contain IS examples, we randomly choose a number of IS intents and treat them as OOS so that the number of examples in these intents are about 25% of the whole training dataset. The full list of chosen intents for each dataset are listed in Appendix A.1.

We reorganize some of the datasets as follows. CLINC150 includes 2 domains, banking and credit card, we evaluate them separately along with the full set. For StackOverflow, 10% examples in training set is stratified split as dev set. For HAR, we remove examples with missing 'answer', and stratified split remaining examples into train, dev, and test set with a 80, 10, and 10 percentage. Different from other selected datasets, ROSTD contains 4,000 OOS examples. ROSTD-coarse is the version that only keep higher hierarchical intent types. Examples in "reminder" intent type from original ROSTD-coarse are treated as ID-OOS. HINT3 consists of 3 domains, including SOFMatress, Curekart and Powerplay11, so we evaluate them separately. 10% of training examples in each of HINT3 datasets is stratified split as dev set.

3.5 Evaluation metrics

Based on current literature, there are 2 types of commonly used metrics for OOS detection.

Threshold dependent metrics are metrics calculated with predicted labels e.g. accuracy. These metrics compare the probability score against a threshold to determine whether a query is considered OOS or not. Also threshold dependent metrics encourage joint evaluation of intent detection and OOS detection that are more suitable under the

Dataset	Train			Dev			Test		
	IS	ID-OOS	OOD-OOS	IS	ID-OOS	OOD-OOS	IS	ID-OOS	OOD-OOS
CLINC150-FULL	11300	3700	100	2260	740	100	3390	1110	1000
CLINC150-BANKING	400	100	0	400	100	600	400	100	1350
CLINC150-CREDIT	400	100	0	400	100	600	400	100	1350
ATIS	4053	425	0	458	42	0	812	81	0
BANKING77	6533	2089	0	1160	380	0	2320	760	0
Stack Overflow	5400	1800	0	600	200	0	6000	2000	0
SNIPS	9371	3713	0	500	200	0	500	200	0
ROSTD	23621	6900	0	3238	943	1500	6661	1960	3090
ROSTD-coarse	23621	6900	0	3238	943	1500	6661	1960	3090
HAR	15893	4592	0	1986	575	0	1985	576	0
HINT3 (SOFmattress)	229	66	0	26	7	0	158	73	166
HINT3 (Powerplay11)	317	102	0	38	14	0	244	31	708
HINT3 (Curekart)	415	125	0	45	15	0	390	62	539

Table 2: **Dataset Statistics.** We preprocess all datasets (details in A.1) and numbers reflect their sizes.

context of VA. Following the literature (Wu et al. (2022), Zhou et al. (2022), Zeng et al. (2021)), the threshold dependent metrics are listed here:

Overall Accuracy is the percentage of examples being correctly classified. For an IS query, it’s predicted correctly if and only if the predicted IS label is correct and the query is predicted as IS. For an OOS query, it should be predicted as OOS to make a correct prediction. **IS Accuracy** is the percentage of correctly predicted IS examples out of all IS examples. **IS F1** is the macro averaged F1 scores across all IS intents. **OOS F1** is the F1 score for OOS examples.

Threshold independent metrics are metrics calculated with a vector of scores each of which measures how confident or likely an OOS algorithm considers a query irrelevant. Such a score is often a number between 0 and 1 where 1 represents IS and 0 represents OOS. This paper follows the literature (Shen et al. (2021), Ryu et al. (2018), Liang et al. (2017), Lee et al. (2018b)) and defines IS as the positive class and OOS as the negative class. We use the metrics for evaluating OOS detection performance: **FPRN**, where N is an integer between 0 and 100, is the false positive rate(FPR) when the true positive rate(TPR) is at least N%. A false positive is an OOS example predicted as IND. We use FPR90 and FPR95. **AUROC** is the area under the Receive Operating Characteristic curve, which measures TPR against FPR at different thresholds. **AUPR_IN** and **AUPR_OUT** are metrics measuring area under the precision-recall curve, when IS and OOS are considered as the positive class, respectively.

<i>Training Example</i>
Can I buy a cell phone ?
<i>Training Entities</i>
entity: cell phone
synonyms: iphone, samsung, galaxy, iphone XR, iphone 11, etc..
<i>Inference Queries</i>
A galaxy is a huge collection of gas, dust, stars and their solar systems.
what is the latest model of galaxy s series?

Table 3: VA designer defines the entity "cell phone" with synonyms. The 1st query contains the word "galaxy" but it is OOD-OOS. The 2nd with "galaxy" is ID-OOS.

3.6 Experiments and Results

We conduct experiment on the benchmark datasets to compare different OOS problem formulations listed in Section 3 (Our discounting approach, Binary Classification (IS/OOS), Multiclass Classification³, and IS Classification utilizing output distribution). As a comparison for the OOS problem formulation only, we keep the IS classification algorithm and OOS algorithm same as our production setup across the 4 formulations, without focusing on the exact choice or implementation of the IS and OOS algorithms. For the discounting method, we use our production intent detection and OOS detection as is. For multi-class classification method, we consider OOS examples as an additional IS intent. For the IS Classification utilizing output distribution formulation, we train an IS classifier with IS training examples and take the max of output confidence vector as the OOS score.

Offline Evaluation Table 4 reports the simple average of metrics across all our benchmark datasets.

³Threshold independent metrics for Multiclass classification is omitted, as our IS classifier outputs confidence vectors (which do not sum up to 1) instead of predicted probability, thus it involves no such component as OOS scores.

Method	Overall Acc.	IS Acc.	IS F1	OOS F1	OOS recall	FPR90	FPR95	AUROC	AUPR_IN	AUPR_OUT
Binary	82.45	86.92	77.87	76.84	74.18	22.76	30.29	91.70	91.21	88.16
Multiclass	74.34	90.36	72.81	74.08	64.84	-	-	-	-	-
IS clf + Max	79.17	85.93	77.99	70.02	65.15	32.66	44.45	87.53	87.15	80.27
Discounting (Our Approach)	84.70	86.43	79.71	80.63	79.31	20.07	27.20	92.64	91.07	89.90

Table 4: **Performance on all datasets** This table compares the discounting method against binary classification, multiclass classification, the IS classifier + max confidence on the full test sets.

Method	Overall Acc.	IS Acc.	IS F1	OOS F1	OOS recall	FPR90	FPR95	AUROC	AUPR_IN	AUPR_OUT
Binary	53.45	72.70	47.97	45.26	38.40	59.24	78.80	81.37	72.07	86.98
Multiclass	52.97	74.65	49.44	56.15	41.15	-	-	-	-	-
IS clf + Max	54.36	75.13	53.46	57.90	43.66	56.38	73.74	76.97	60.33	86.01
Discounting (Our Approach)	61.46	69.38	51.79	60.05	54.94	50.84	71.02	82.83	69.32	88.97

Table 5: **Performance on HINT3** This table compares the discounting method against the Multiclass classification method, the binary classification method, the IS classifier + max confidence on the full test sets.

The discounting approach achieves better performance across most metrics. We report the average metrics in Table 5 on the subset of the 3 HINT3 datasets, which are designed to represent real world imbalanced datasets.

Table 6 compares the Multi-Class formulation against our formulation on all datasets. Our approach performs on par on ID-OOS but generalize better to OOD-OOS. In real world application, limited ID-OOS is provided by customer during training but the algorithm is expected to perform well on both categories without overfitting.

Method	Test set	Overall Acc.	IS Acc.	IS F1	OOS F1	OOS recall
K+1 Classes	IND+ID-OOS	90.65	90.36	86.46	90.82	92.76
discounting	IND+ID-OOS	86.14	88.04	84.62	76.25	80.81
K+1 Classes	IND+OOD-OOS	60.28	89.41	65.07	46.60	32.69
discounting	IND+OOD-OOS	78.31	85.72	76.71	74.36	71.43
K+1 Classes	IND+both OOS types	74.34	90.36	72.81	74.08	64.84
discounting	IND+both OOS types	84.70	86.43	79.71	80.63	79.31

Table 6: **Performance on various test sets** We compare the discounting method against the Multiclass classification method on 3 versions of test sets: IS + ID-OOS, IS + OOD-OOS (average across the datasets with OOD-OOS test examples), IS + both types of OOS examples.

Online Evaluation During real-world deployment of this algorithm, we conducted additional online evaluation by analyzing the output distribution change on real production traffic because chatbot designers typically rely on output confidence scores to make business decisions eg. jumping to a node in the dialog tree, handing off to human agents or asking a follow-up question. Therefore, we deployed the proposed OOS algorithm in production and monitored different statistics on 10% of randomly selected real traffic for months before surfacing it to end-users. We observed that more than 85% of live traffic queries will have a less than 10% change in top confidence after the change in OOS algorithms

(Full distribution shown in Figure 1 in Appendix). For enterprise customers with complex dialog conditions, a new algorithm that does not disrupt the normal workflow is critical for adoption.

Computational Efficiency and Scalability Our product has a training set size limit of 25k IS and OOS training examples each and 2k IS classes. Based on this maximum training set size setting, the maximum model size for OOS detection is less 150MB based on offline testing. Based on online testing, the 99 percentiles for training time and model size of our OOS algorithm are within 30 seconds and 70MB, respectively.

4 Conclusion

The paper presents a novel Out of Scope (OOS) detection component in a task-oriented dialog system. It allows the assistant to recognize user input that is not designed to be answered by the assistant and need to be handed off to a human agent. For business, a well designed Out of Scope detection system can improve customer satisfaction, user engagement, lead generation and saves cost. On one hand, business wants the assistant to hand off quickly when a user input is Out of Scope. On the other hand, unnecessary hand off could increase human intervention and reduce the value of VA. We design an OOS detection system that overcomes a multitude of real-world challenges, and deploy it in production.⁴ We list out the lessons learned and both offline and online evaluation techniques for designing a robust, efficient and scalable system for enterprise VA platform.

⁴<https://cloud.ibm.com/docs/assistant?topic=assistant-irrelevance-detection>,
<https://cloud.ibm.com/docs/assistant?topic=assistant-release-notes#assistant-jun162022>

Limitations

Extensive benchmarking for other languages is out-of-scope for this work, but we have extended the approach to many European languages in the product with similar gains in performance (Wang et al., 2022). Code switching isn’t evaluated in this work, but it is commonly observed in chatbots deployed in the wild.

We have not discussed synthetic OOS examples. Despite its demonstrated effectiveness, they need caution in real world production system from a robustness perspective: it’s possible to introduce spurious correlation by generated synthetic data.

References

- Gaurav Arora, Chirag Jain, Manas Chaturvedi, and Krupal Modi. 2020. Hint3: Raising the bar for intent detection in the wild. *arXiv preprint arXiv:2009.13833*.
- Iñigo Casanueva, Tadas Temcinas, Daniela Gerz, Matthew Henderson, and Ivan Vulic. 2020. [Efficient intent detection with dual sentence encoders](#). In *Proceedings of the 2nd Workshop on NLP for ConvAI - ACL 2020*. Data available at <https://github.com/PolyAI-LDN/task-specific-datasets>.
- DongHyun Choi, Myeongcheol Shin, EungGyun Kim, and Dong Ryeol Shin. 2021. [Outflip: Generating out-of-domain samples for unknown intent detection with natural language attack](#). *CoRR*, abs/2105.05601.
- Alice Coucke, Alaa Saade, Adrien Ball, Théodore Bluche, Alexandre Caulier, David Leroy, Clément Doumouro, Thibault Gisselbrecht, Francesco Caltagirone, Thibaut Lavril, Maël Primet, and Joseph Dureau. 2018. [Snips voice platform: an embedded spoken language understanding system for private-by-design voice interfaces](#). *CoRR*, abs/1805.10190.
- Varun Gangal, Abhinav Arora, Arash Einolghozati, and Sonal Gupta. 2020. Likelihood ratios and generative classifiers for unsupervised out-of-domain detection in task oriented dialog. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 34, pages 7764–7771.
- Charles T Hemphill, John J Godfrey, and George R Doddington. 1990. The atis spoken language systems pilot corpus. In *Speech and Natural Language: Proceedings of a Workshop Held at Hidden Valley, Pennsylvania, June 24-27, 1990*.
- Dan Hendrycks and Kevin Gimpel. 2016. [A baseline for detecting misclassified and out-of-distribution examples in neural networks](#). *CoRR*, abs/1610.02136.
- Amita Kamath, Robin Jia, and Percy Liang. 2020. [Selective question answering under domain shift](#). *CoRR*, abs/2006.09462.
- Stefan Larson, Anish Mahendran, Joseph J Peper, Christopher Clarke, Andrew Lee, Parker Hill, Jonathan K Kummerfeld, Kevin Leach, Michael A Laurenzano, Lingjia Tang, et al. 2019. An evaluation dataset for intent classification and out-of-scope prediction. *arXiv preprint arXiv:1909.02027*.
- Kimin Lee, Honglak Lee, Kibok Lee, and Jinwoo Shin. 2018a. [Training confidence-calibrated classifiers for detecting out-of-distribution samples](#). In *International Conference on Learning Representations*.
- Kimin Lee, Kibok Lee, Honglak Lee, and Jinwoo Shin. 2018b. [A simple unified framework for detecting out-of-distribution samples and adversarial attacks](#). In *Advances in Neural Information Processing Systems*, volume 31. Curran Associates, Inc.
- David D. Lewis and William A. Gale. 1994. A sequential algorithm for training text classifiers. In *SIGIR ’94*, pages 3–12, London. Springer London.
- Shiyu Liang, Yixuan Li, and R. Srikant. 2017. [Principled detection of out-of-distribution examples in neural networks](#). *CoRR*, abs/1706.02690.
- Ting-En Lin and Hua Xu. 2019. [Deep unknown intent detection with margin loss](#). *CoRR*, abs/1906.00434.
- Xingkun Liu, Arash Eshghi, Pawel Swietojanski, and Verena Rieser. 2019. Benchmarking natural language understanding services for building conversational agents. *arXiv preprint arXiv:1903.05566*.
- Haode Qi, Lin Pan, Atin Sood, Abhishek Shah, Ladislav Kunc, and Saloni Potdar. 2020. [Benchmarking intent detection for task-oriented dialog systems](#). *CoRR*, abs/2012.03929.
- Seonghan Ryu, Sangjun Koo, Hwanjo Yu, and Gary Geunbae Lee. 2018. [Out-of-domain detection based on generative adversarial network](#). In *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing*, pages 714–718, Brussels, Belgium. Association for Computational Linguistics.
- Yilin Shen, Yen-Chang Hsu, Avik Ray, and Hongxia Jin. 2021. [Enhancing the generalization for intent classification and out-of-domain detection in SLU](#). *CoRR*, abs/2106.14464.
- David M. J. Tax and Robert P. W. Duin. 1999. [Data domain description using support vectors](#). In *ESANN 1999, 7th European Symposium on Artificial Neural Networks, Bruges, Belgium, April 21-23, 1999, Proceedings*, pages 251–256.
- Gengyu Wang, Cheng Qian, Lin Pan, Haode Qi, Ladislav Kunc, and Saloni Potdar. 2022. Benchmarking language-agnostic intent classification for virtual assistant platforms. In *Proceedings of the Workshop on Multilingual Information Access (MIA)*, pages 69–76.

Yanan Wu, Keqing He, Yuanmeng Yan, QiXiang Gao, Zhiyuan Zeng, Fujia Zheng, Lulu Zhao, Huixing Jiang, Wei Wu, and Weiran Xu. 2022. [Revisit overconfidence for OOD detection: Reassigned contrastive learning with adaptive class-dependent threshold](#). In *Proceedings of the 2022 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 4165–4179, Seattle, United States. Association for Computational Linguistics.

Eyup Halit Yilmaz and Cagri Toraman. 2022. [D2U: distance-to-uniform learning for out-of-scope detection](#). In *Proceedings of the 2022 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, NAACL 2022, Seattle, WA, United States, July 10-15, 2022*, pages 2093–2108. Association for Computational Linguistics.

Zhiyuan Zeng, Keqing He, Yuanmeng Yan, Zijun Liu, Yanan Wu, Hong Xu, Huixing Jiang, and Weiran Xu. 2021. [Modeling discriminative representations for out-of-domain detection with supervised contrastive learning](#). *CoRR*, abs/2105.14289.

Li-Ming Zhan, Haowen Liang, Bo Liu, Lu Fan, Xiaoming Wu, and Albert Y. S. Lam. 2021. [Out-of-scope intent detection with self-supervision and discriminative training](#). *CoRR*, abs/2106.08616.

Jianguo Zhang, Kazuma Hashimoto, Yao Wan, Zhiwei Liu, Ye Liu, Caiming Xiong, and Philip Yu. 2022. [Are pre-trained transformers robust in intent classification? a missing ingredient in evaluation of out-of-scope intent detection](#). In *Proceedings of the 4th Workshop on NLP for Conversational AI*, pages 12–20, Dublin, Ireland. Association for Computational Linguistics.

Yunhua Zhou, Peiju Liu, and Xipeng Qiu. 2022. [KNN-contrastive learning for out-of-domain intent classification](#). In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 5129–5141, Dublin, Ireland. Association for Computational Linguistics.

A Appendix

A.1 List of IN-OOS intents

Here we list the intents for each dataset that are treated as IN-OOS intents in our benchmark.

Stackoverflow: python

SNIPS: SearchCreativeWork and SearchScreeningEvent

HAR: intents, hue_lightoff, explain, remove, addcontact, wemo_on, podcasts, createoradd, music, praise, radio, dontcare

ROSTD: reminder/set_reminder, reminder/cancel_reminder, reminder/show_reminders

HINT3 SOFMattress:
SIZE_CUSTOMIZATION,
ABOUT_SOF_MATTRESS, LEAD_GEN,
COMPARISON, WARRANTY, DELAY_IN_DELIVERY

HINT3 Powerplay11:
NO_EMAIL_CONFIRMATION,
TEAM_DEADLINE, FAKE_TEAMS,
CANNOT_SEE_JOINED_CONTESTS, REFUND_OF_ADDED_CASH, HOW_TO_PLAY,
FEEDBACK, ACCOUNT_NOT_VERIFIED,
DEDUCTED_AMOUNT_NOT_RECEIVED,
CRITICISM, NEW_TEAM_PATTERN, OFFERS_AND_REFERRALS

HINT3 Curekart: EXPIRY_DATE,
CONSULT_START, CHECK_PINCODE,
ORDER_TAKING, INTERNATIONAL_SHIPPING, IMMUNITY,
SIDE_EFFECT, START_OVER, PORTAL_ISSUE, MODES_OF_PAYMENTS, ORDER_QUERY, SIGN_UP, WORK_FROM_HOME

A.2 Our OOS problem formulation is algorithm-agnostic:

We conducted the same experiment with another OOS algorithm: autoencoder with reconstruction loss as OOS score. The findings are similar: our OOS formulation demonstrate advantages over others. Detailed metrics are shown in Table 7.

A.3 Online evaluation statistics

Figure 1 shows the full distribution of differences in top confidence between the proposed OOS algorithms vs previous OOS algorithm on a percentage of live traffic

Distribution of difference in top confidences between 2 OOS algorithm:




Figure 1: Distribution of differences in top confidence between the proposed OOS algorithms vs previous OOS algorithm on a percentage of live traffic

Method	Overall Acc.	IS Acc.	IS F1	OOS F1	OOS recall	FPR90	FPR95	AUROC	AUPR_IN	AUPR_OUT
Binary	83.25	85.09	76.68	81.70	77.16	28.24	34.78	91.47	90.37	89.27
Multiclass	74.34	90.36	72.81	74.08	64.84					
IS clf + Max	79.17	85.93	77.99	70.02	65.15	32.66	44.45	87.53	87.15	80.27
Discounting (Our Approach)	84.30	85.85	77.88	83.17	79.35	19.65	25.56	93.18	92.24	91.63

Table 7: **Performance metrics** This table compares the discounting method against the Multiclass classification method, the binary classification method, the IS classifier + max confidence on the full test sets, using autoencoder as the OOS detection algorithm.